

## Assessing Balance In HTE ~ ICLR DB

```
rm(list=ls())
```

```
library(ggplot2)
library(tidyr)
library(psych)
```

```
##
## Attaching package: 'psych'

## The following objects are masked from 'package:ggplot2':
##
##      %+%, alpha
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##      filter, lag

## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
```

```
library(grf)
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --

## v tibble  3.1.8      v stringr 1.4.1
## v readr   2.1.3      v forcats 0.5.2
## v purrr    0.3.5

## -- Conflicts ----- tidyverse_conflicts() --
## x psych::%+%( ) masks ggplot2::%+%( )
## x psych::alpha( ) masks ggplot2::alpha( )
## x dplyr::filter( ) masks stats::filter( )
## x dplyr::lag( ) masks stats::lag( )
```

```
require(gridExtra)
```

```
## Loading required package: gridExtra
##
## Attaching package: 'gridExtra'
##
## The following object is masked from 'package:dplyr':
##
##      combine
```

```
library(xtable)
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##      select
```

```
library(sandwich)
```

```
# Auxiliary function to computes adjusted p-values
# following the Romano-Wolf method.
# For a reference, see http://ftp.iza.org/dp12845.pdf page 8
# t.orig: vector of t-statistics from original model
# t.boot: matrix of t-statistics from bootstrapped models
romano_wolf_correction <- function(t.orig, t.boot) {
  abs.t.orig <- abs(t.orig)
  abs.t.boot <- abs(t.boot)
  abs.t.sorted <- sort(abs.t.orig, decreasing = TRUE)

  max.order <- order(abs.t.orig, decreasing = TRUE)
  rev.order <- order(max.order)

  M <- nrow(t.boot)
  S <- ncol(t.boot)

  p.adj <- rep(0, S)
  p.adj[1] <- mean(apply(abs.t.boot, 1, max) > abs.t.sorted[1])
  for (s in seq(2, S)) {
    cur.index <- max.order[s:S]
    p.init <- mean(apply(abs.t.boot[, cur.index, drop=FALSE], 1, max) > abs.t.sorted[s])
    p.adj[s] <- max(p.init, p.adj[s-1])
  }
}
```

```

    p.adj[rev.order]
  }

# Computes adjusted p-values for linear regression (lm) models.
# model: object of lm class (i.e., a linear reg model)
# indices: vector of integers for the coefficients that will be tested
# cov.type: type of standard error (to be passed to sandwich::vcovHC)
# num.boot: number of null bootstrap samples. Increase to stabilize across runs.
# Note: results are probabilistic and may change slightly at every run.
#
# Adapted from the p_adjust from the hdm package, written by Philipp Bach.
# https://github.com/PhilippBach/hdm_prev/blob/master/R/p_adjust.R
summary_rw_lm <- function(model, indices=NULL, cov.type="HC2", num.boot=10000) {

  if (is.null(indices)) {
    indices <- 1:nrow(coef(summary(model)))
  }
  # Grab the original t values.
  summary <- coef(summary(model))[indices,,drop=FALSE]
  t.orig <- summary[, "t value"]

  # Null resampling.
  # This is a trick to speed up bootstrapping linear models.
  # Here, we don't really need to re-fit linear regressions, which would be a bit slow.
  # We know that betahat ~ N(beta, Sigma), and we have an estimate Sigma.hat.
  # So we can approximate "null t-values" by
  # - Draw beta.boot ~ N(0, Sigma.hat) --- note the 0 here, this is what makes it a *null* t-value.
  # - Compute t.boot = beta.boot / sqrt(diag(Sigma.hat))
  Sigma.hat <- vcovHC(model, type=cov.type)[indices, indices]
  se.orig <- sqrt(diag(Sigma.hat))
  num.coef <- length(se.orig)
  beta.boot <- mvrnorm(n=num.boot, mu=rep(0, num.coef), Sigma=Sigma.hat)
  t.boot <- sweep(beta.boot, 2, se.orig, "/")
  p.adj <- romano_wolf_correction(t.orig, t.boot)

  result <- cbind(summary[,c(1,2,4),drop=F], p.adj)
  colnames(result) <- c('Estimate', 'Std. Error', 'Orig. p-value', 'Adj. p-value')
  result
}

```

```
setwd('~/.Mirror/github/aICLR/')
```

```
df_authors = read.csv("./data/database/outputs/df_authors.csv")
df_prestige = read.csv("./data/database/outputs/df_prestige.csv")
```

```

Authors_2017 = df_authors %>% group_by(author_id) %>%
  filter(conf_year == 2017, author_no == 1, current_position_flag == 1)

Authors_2018 = df_authors %>% group_by(author_id) %>%
  filter(conf_year == 2018, author_no == 1, current_position_flag == 1)

Authors_2019 = df_authors %>% group_by(author_id) %>%
  filter(conf_year == 2019, author_no == 1, current_position_flag == 1)

```

```

g1 = ggplot(Authors_2017, aes(US_Canada)) +
  geom_bar() +
  ggtitle("2017 1st Author US Canada") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g2 = ggplot(Authors_2018, aes(US_Canada)) +
  geom_bar() +
  ggtitle("2018 1st Author US Canada") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g3 = ggplot(Authors_2019, aes(US_Canada)) +
  geom_bar() +
  ggtitle("2019 1st Author US Canada") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")
grid.arrange(g1, g2, g3, ncol=2)

```

```

## Warning: 'stat(count)' was deprecated in ggplot2 3.4.0.
## i Please use 'after_stat(count)' instead.

```

```

## Warning: Removed 118 rows containing non-finite values ('stat_count()').
## Removed 118 rows containing non-finite values ('stat_count()').

```

```

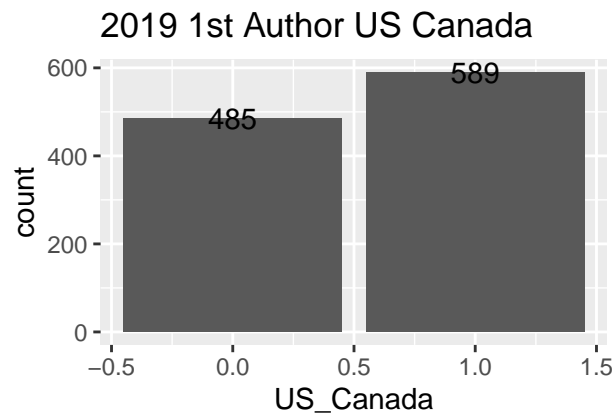

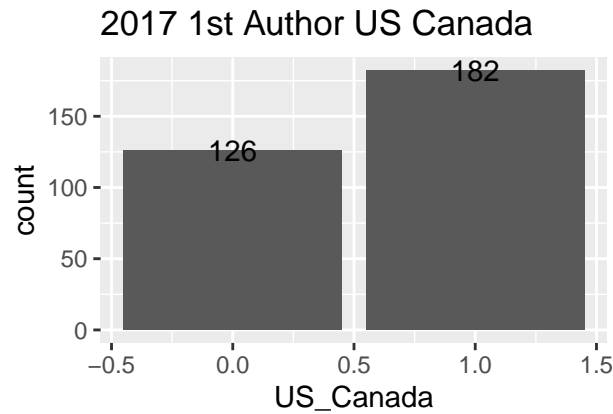
## Warning: Removed 221 rows containing non-finite values ('stat_count()').
## Removed 221 rows containing non-finite values ('stat_count()').

```

```

## Warning: Removed 228 rows containing non-finite values ('stat_count()').
## Removed 228 rows containing non-finite values ('stat_count()').

```




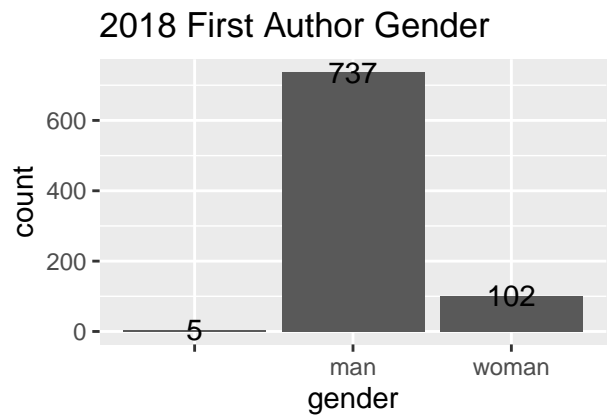
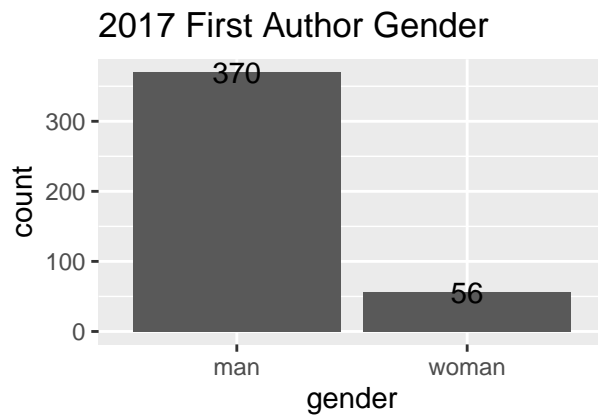
```
Authors_2017 = df_authors %>% filter(conf_year == 2017, author_no == 1, current_position_flag == 1)
Authors_2018 = df_authors %>% filter(conf_year == 2018, author_no == 1, current_position_flag == 1)
Authors_2019 = df_authors %>% filter(conf_year == 2019, author_no == 1, current_position_flag == 1)

g1 = ggplot(Authors_2017, aes(gender)) +
  geom_bar() +
  ggtitle("2017 First Author Gender") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g2 = ggplot(Authors_2018, aes(gender)) +
  geom_bar() +
  ggtitle("2018 First Author Gender") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g3 = ggplot(Authors_2019, aes(gender)) +
  geom_bar() +
  ggtitle("2019 First Author Gender") +
  stat_count(geom = "text",
    aes(label = stat(count)),
```

```
position="stack", colour="black")
grid.arrange(g1, g2, g3, ncol=2)
```


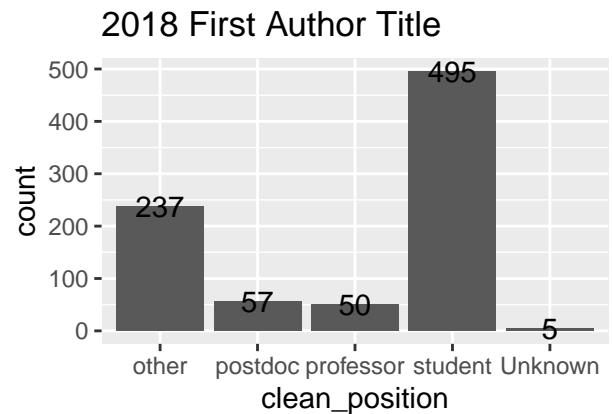
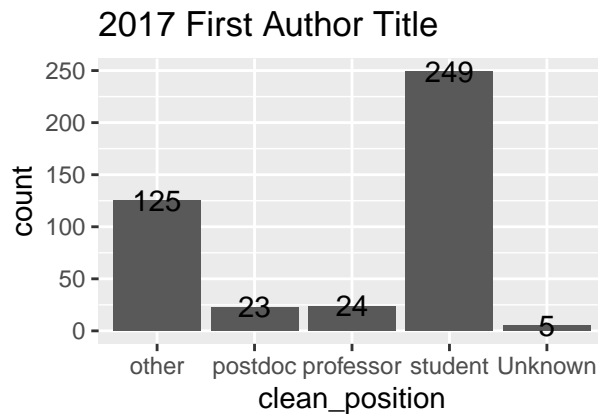


```
Authors_2017 = df_authors %>% filter(conf_year == 2017, author_no == 1, current_position_flag == 1)
Authors_2018 = df_authors %>% filter(conf_year == 2018, author_no == 1, current_position_flag == 1)
Authors_2019 = df_authors %>% filter(conf_year == 2019, author_no == 1, current_position_flag == 1)

g1 = ggplot(Authors_2017, aes(clean_position)) +
  geom_bar() +
  ggtitle("2017 First Author Title") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g2 = ggplot(Authors_2018, aes(clean_position)) +
  geom_bar() +
  ggtitle("2018 First Author Title") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")
```

```
g3 = ggplot(Authors_2019, aes(clean_position)) +
  geom_bar() +
  ggtitle("2019 First Author Title") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")
grid.arrange(g1, g2, g3, ncol=2)
```



```
Authors_2017 = df_authors %>% group_by(submission_id) %>%
  top_n(1, author_no) %>%
  filter(conf_year == 2017, current_position_flag == 1)
```

```
Authors_2018 = df_authors %>% group_by(submission_id) %>%
  top_n(1, author_no) %>%
  filter(conf_year == 2018, current_position_flag == 1)
```

```
Authors_2019 = df_authors %>%
  group_by(submission_id) %>%
  top_n(1, author_no) %>%
  filter(conf_year == 2019, current_position_flag == 1)
```

```
g1 = ggplot(Authors_2017, aes(clean_position)) +
  geom_bar() +
  ggtitle("2017 Last Author Title") +
```

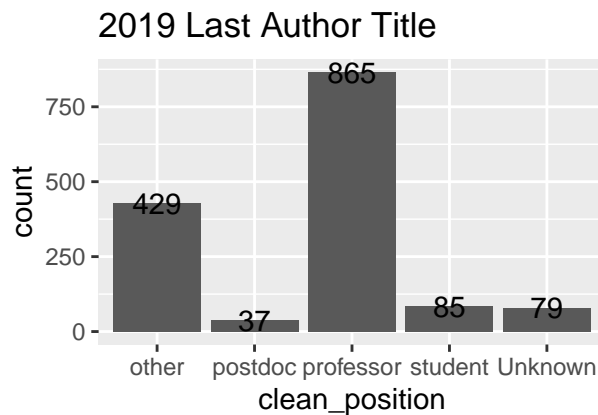

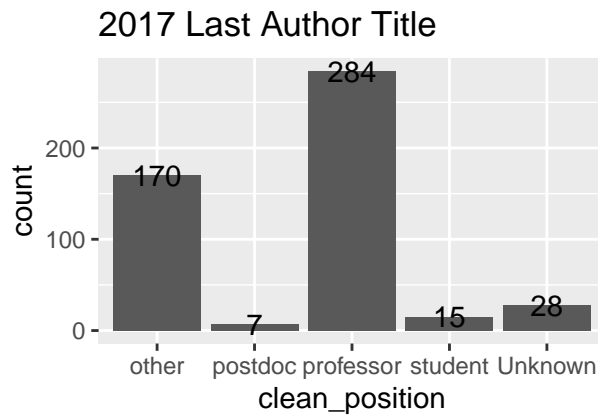
```

stat_count(geom = "text",
  aes(label = stat(count)),
  position="stack", colour="black")

g2 = ggplot(Authors_2018, aes(clean_position)) +
  geom_bar() +
  ggtitle("2018 Last Author Title") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g3 = ggplot(Authors_2019, aes(clean_position)) +
  geom_bar() +
  ggtitle("2019 Last Author Title") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")
grid.arrange(g1, g2, g3, ncol=2)

```



```

df_prestige = df_prestige %>% mutate(MAX_CITE_Ntile = ntile(MAX_CITE, 4))
df_prestige_2017 = df_prestige %>%
  filter(conf_year == 2017)

df_prestige_2018 = df_prestige %>%

```



```

filter(conf_year ==2018)

df_prestige_2019 = df_prestige %>%
  filter(conf_year ==2019)

g1 = ggplot(df_prestige_2017, aes(MAX_CITE_Ntile)) +
  geom_bar() +
  ggtitle("2017 Last Author Citation Precentile \n Group") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g2 = ggplot(df_prestige_2018, aes(MAX_CITE_Ntile)) +
  geom_bar() +
  ggtitle("2018 Last Author Citation Precentile \n Group") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")

g3 = ggplot(df_prestige_2019, aes(MAX_CITE_Ntile)) +
  geom_bar() +
  ggtitle("2019 Last Author Citation Precentile \n Group") +
  stat_count(geom = "text",
    aes(label = stat(count)),
    position="stack", colour="black")
grid.arrange(g1, g2, g3,ncol=2)

```

```

## Warning: Removed 45 rows containing non-finite values ('stat_count()').
## Removed 45 rows containing non-finite values ('stat_count()').

```

```




## Warning: Removed 114 rows containing non-finite values ('stat_count()').
## Removed 114 rows containing non-finite values ('stat_count()').

```

```

## Warning: Removed 148 rows containing non-finite values ('stat_count()').
## Removed 148 rows containing non-finite values ('stat_count()').

```



```
features = c("author_no","conf_year.x","US_Canada","clean_position","gender","AVG_len" , "AVG_confidence")
df_combined = merge(df_authors, df_prestige,by.x = 'submission_id', by.y = 'id') %>% filter(current_position_flag==1)

df_combined_limited_features = df_combined[features]
```


```
df_authors %>% filter(author_no ==1 & current_position_flag==1) %>% group_by(conf_year) %>% summarize(n = n())
```

```
## # A tibble: 3 x 2
##   conf_year      n
##   <int> <int>
## 1    2017   426
## 2    2018   844
## 3    2019  1302
```

```
ggcorr(df_combined_limited_features,hjust = 0.5,vjust=.2, size = 3) + ggplot2::labs(title = "Pearson Correlation Matrix") +
  theme(plot.title = element_text(hjust = 0.5))
```

```
## Warning in ggcorr(df_combined_limited_features, hjust = 0.5, vjust = 0.2, : data
## in column(s) 'clean_position', 'gender' are not numeric and were ignored
```

## Pearson Correlation Matrix



```
df_combined_limited_features_first_author = df_combined_limited_features %>% filter(author_no == 1 & conf_year.x == 2014)

# Fit the full model
full.model <- lm(polite_prediction ~ ., data = na.omit(df_combined_limited_features_first_author))
library(MASS)
step.model <- stepAIC(full.model, direction = "both",
                      trace = FALSE)
summary(step.model)
```

```
##
## Call:
## lm(formula = polite_prediction ~ conf_year.x + AVG_len + MAX_CITE_Ntile +
##     AVG_rating + MAX_CITE + AVG_CITE, data = na.omit(df_combined_limited_features_first_author))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.43707 -0.21804  0.01914  0.11689  0.81997
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.288e+02  3.618e+01   3.561 0.000390 ***
## conf_year.x    -6.386e-02  1.793e-02  -3.561 0.000390 ***
## AVG_len         1.810e-04  4.638e-05   3.902 0.000103 ***
## MAX_CITE_Ntile -1.483e-02  9.104e-03  -1.629 0.103649
## AVG_rating      4.601e-02  6.911e-03   6.658 5e-11 ***
```

```
## MAX_CITE      3.373e-06  1.535e-06   2.197 0.028261 *
## AVG_CITE     -1.029e-05  4.930e-06  -2.086 0.037260 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2362 on 841 degrees of freedom
## Multiple R-squared:  0.08147,    Adjusted R-squared:  0.07492
## F-statistic: 12.43 on 6 and 841 DF,  p-value: 1.882e-13

full.model <- lm(polite_prediction ~., data = na.omit(df_combined_limited_features_first_author))
summary(full.model)
```

```
##
## Call:
## lm(formula = polite_prediction ~ ., data = na.omit(df_combined_limited_features_first_author))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.42777 -0.21542  0.01786  0.11429  0.82106
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.285e+02  3.645e+01   3.525 0.000447 ***
## author_no             NA         NA      NA      NA
## conf_year.x    -6.375e-02  1.806e-02  -3.530 0.000439 ***
## US_Canada        1.637e-02  1.732e-02   0.945 0.344882
## clean_positionpostdoc  5.581e-02  4.654e-02   1.199 0.230821
## clean_positionprofessor 9.041e-02  4.569e-02   1.979 0.048190 *
## clean_positionstudent  4.268e-02  3.625e-02   1.177 0.239340
## clean_positionUnknown  3.723e-02  8.347e-02   0.446 0.655664
## genderman         4.037e-02  1.069e-01   0.378 0.705725
## genderwoman       4.432e-02  1.089e-01   0.407 0.684049
## AVG_len           1.754e-04  4.678e-05   3.750 0.000189 ***
## AVG_confidence     8.023e-03  1.621e-02   0.495 0.620797
## MAX_CITE_Ntile    -1.602e-02  9.301e-03  -1.722 0.085402 .
## AVG_rating        4.620e-02  7.133e-03   6.477 1.6e-10 ***
## MAX_CITE          3.443e-06  1.540e-06   2.235 0.025666 *
## AVG_CITE         -1.022e-05  4.952e-06  -2.063 0.039413 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2365 on 833 degrees of freedom
## Multiple R-squared:  0.08802,    Adjusted R-squared:  0.07269
## F-statistic: 5.743 on 14 and 833 DF,  p-value: 8.95e-11
```

```
full.model <- lm(AVG_rating ~. + conf_year.x*MAX_CITE_Ntile, data = na.omit(df_combined_limited_features))
summary(full.model)

##
## Call:
## lm(formula = AVG_rating ~ . + conf_year.x * MAX_CITE_Ntile, data = na.omit(df_combined_limited_features))
##
## Residuals:
```

```
##      Min      1Q  Median      3Q      Max
## -2.9514 -0.8060 -0.0361  0.7773  3.4613
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -1.871e+02  4.015e+02  -0.466   0.6414
## author_no              NA         NA      NA      NA
## conf_year.x      9.585e-02  1.990e-01   0.482   0.6302
## US_Canada       1.401e-01  8.205e-02   1.708   0.0881 .
## clean_positionpostdoc -7.634e-02  2.211e-01  -0.345   0.7300
## clean_positionprofessor -6.766e-02  2.172e-01  -0.311   0.7555
## clean_positionstudent -1.325e-02  1.720e-01  -0.077   0.9386
## clean_positionUnknown  5.793e-01  3.954e-01   1.465   0.1433
## genderman        2.166e-01  5.068e-01   0.427   0.6693
## genderwoman      1.008e-02  5.164e-01   0.020   0.9844
## AVG_len         -1.711e-04  2.236e-04  -0.765   0.4443
## AVG_confidence   -4.694e-01  7.532e-02  -6.232 7.29e-10 ***
## polite_prediction  1.032e+00  1.605e-01   6.431 2.14e-10 ***
## MAX_CITE_Ntile    1.145e+02  1.491e+02   0.768   0.4429
## MAX_CITE          4.842e-06  7.341e-06   0.660   0.5097
## AVG_CITE         -4.493e-05  2.349e-05  -1.913   0.0561 .
## conf_year.x:MAX_CITE_Ntile -5.661e-02  7.391e-02  -0.766   0.4439
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.121 on 832 degrees of freedom
## Multiple R-squared:  0.1463, Adjusted R-squared:  0.131
## F-statistic: 9.509 on 15 and 832 DF,  p-value: < 2.2e-16
```

```
df_submissions_1718 = df_combined_limited_features_first_author %>% filter(conf_year.x %in% c(2017,2018))
df_submissions_1718 = df_submissions_1718 %>% mutate(W = if_else(conf_year.x==2017,0,1))

covariates = c("US_Canada","clean_position","gender","AVG_len" ,"AVG_confidence","MAX_CITE","polite_prediction")
#covariates = c("AVG_len" ,"MAX_CITE_Ntile","polite_prediction","AVG_confidence")

df_submissions_1718 = na.omit(df_submissions_1718)
XX <- model.matrix(formula(paste0("~", paste0(covariates, collapse="+"))), data=df_submissions_1718)

set.seed(1)


forest <- causal_forest(
  X=XX,
  W=df_submissions_1718[,"W"],
  Y=df_submissions_1718[,"AVG_rating"]
  #,num.trees = 100
)

forest.ate <- average_treatment_effect(forest)

forest.ate
```

```
##      estimate      std.err
## -0.06920760  0.09428644
```

```
hist(forest$W.hat, main="Estimated propensity scores \n(causal forest with submission data)", xlim=c(-.1, 1.1))
```



```
covariates = c("US_Canada", "AVG_len", "AVG_confidence", "MAX_CITE", "polite_prediction", "gender", "clean_p
# Here, adding covariates and their interactions, though there are many other possibilities.
fmla <- formula(paste("~ 0 +", paste(apply(expand.grid(covariates, covariates), 1, function(x) paste0(x
# Using the propensity score estimated above
#check to see if you are using ATE causal forest ICLR_analysis_Submission_1718
e.hat <- forest$W.hat

XX <- model.matrix(fmla, df_submissions_1718)
W <- df_submissions_1718[, "W"]
pp <- ncol(XX)

# Unadjusted covariate means, variances and standardized abs mean differences
means.treat <- apply(XX[W == 1,], 2, mean)
means.ctrl <- apply(XX[W == 0,], 2, mean)
abs.mean.diff <- abs(means.treat - means.ctrl)

var.treat <- apply(XX[W == 1,], 2, var)
var.ctrl <- apply(XX[W == 0,], 2, var)
std <- sqrt(var.treat + var.ctrl)
```

```


# Adjusted covariate means, variances and standardized abs mean differences
means.treat.adj <- apply(XX*W/e.hat, 2, mean)
means.ctrl.adj <- apply(XX*(1-W)/(1-e.hat), 2, mean)
abs.mean.diff.adj <- abs(means.treat.adj - means.ctrl.adj)

var.treat.adj <- apply(XX*W/e.hat, 2, var)
var.ctrl.adj <- apply(XX*(1-W)/(1-e.hat), 2, var)
std.adj <- sqrt(var.treat.adj + var.ctrl.adj)

# Plotting
#png(file = "balance_ICLR_analysis_Submission_1718_ATE.png")
par(oma=c(0,4,0,0))
plot(-2, xaxt="n", yaxt="n", xlab="", ylab="", xlim=c(-.01, 1.3), ylim=c(0, pp+1), main="Standardized absolute mean differences",
axis(side=1, at=c(-1, 0, 1), las=1)
lines(abs.mean.diff / std, seq(1, pp), type="p", col="blue", pch=19)
lines(abs.mean.diff.adj / std.adj, seq(1, pp), type="p", col="orange", pch=19)
legend("topright", c("Unadjusted", "Adjusted"), col=c("blue", "orange"), pch=19)
abline(v = seq(0, 1, by=.25), lty = 2, col = "grey", lwd=.5)
abline(h = 1:pp, lty = 2, col = "grey", lwd=.5)
mtext(colnames(XX), side=2, cex=0.42, at=1:pp, padj=.4, adj=1, col="black", las=1, line=.3)
abline(v = 0)

```

## Standardized absolute mean differences Submission Level Data 2017–2018



```

set.seed(1)
group = "US_Canada"

```

```

covariates = c("US_Canada", "clean_position", "gender", "AVG_len", "AVG_confidence", "MAX_CITE", "polite_pre

XX <- model.matrix(formula(paste0("~", paste0(covariates, collapse="+"))), data=df_submissions_1718)
W=df_submissions_1718[, "W"]
Y=df_submissions_1718[, "AVG_rating"]

forest.tau <- causal_forest(XX, Y, W)

tau.hat <- predict(forest.tau)$predictions
m.hat <- forest.tau$Y.hat #  $E[Y|X]$  estimates
e.hat <- forest.tau$W.hat #  $e(X) := E[W|X]$  estimates (or known quantity)
tau.hat <- forest.tau$predictions #  $\tau(X)$  estimates

# Predicting  $\mu.hat(X[i], 1)$  and  $\mu.hat(X[i], 0)$  for obs in held-out sample
# Note: to understand this, read equations 6-8 in this vignette
# https://grf-labs.github.io/grf/articles/muhats.html
mu.hat.0 <- m.hat - e.hat * tau.hat #  $E[Y|X, W=0] = E[Y|X] - e(X) * \tau(X)$ 
mu.hat.1 <- m.hat + (1 - e.hat) * tau.hat #  $E[Y|X, W=1] = E[Y|X] + (1 - e(X)) * \tau(X)$ 

# Compute AIPW scores
aipw.scores <- tau.hat + W / e.hat * (Y - mu.hat.1) - (1 - W) / (1 - e.hat) * (Y - mu.hat.0)

# Estimate average treatment effect conditional on group membership
fmla <- formula(paste0('aipw.scores ~ factor(', group, ')'))
ols <- lm(fmla, data=transform(df_submissions_1718[covariates], aipw.scores=aipw.scores))

summary_rw_lm(ols)

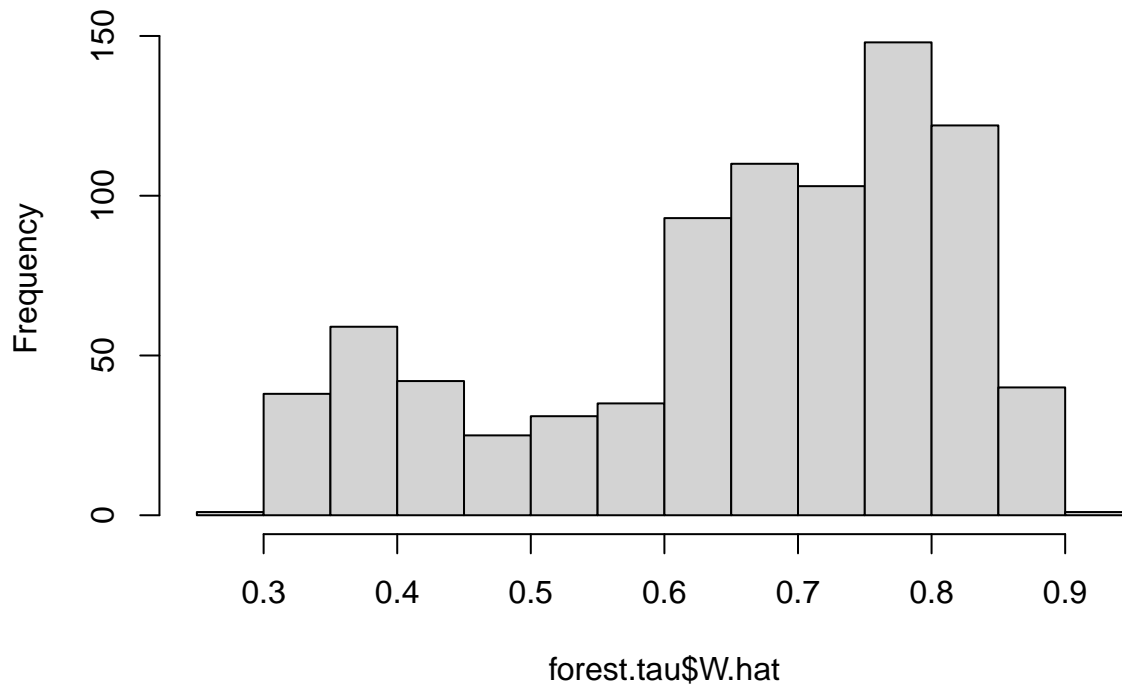
##               Estimate Std. Error Orig. p-value Adj. p-value
## (Intercept)    0.0271899  0.1493715     0.8556038     0.8574
## factor(US_Canada)1 -0.1602845  0.1926107     0.4055490     0.5562

hist(forest.tau$W.hat)

```



## Histogram of forest.tau\$W.hat



```
test_calibration(forest.tau)
```

```
##
## Best linear fit using forest predictions (on held-out data)
## as well as the mean forest prediction as regressors, along
## with one-sided heteroskedasticity-robust (HC3) SEs:
##
##               Estimate Std. Error t value Pr(>t)
## mean.forest.prediction    1.0499    1.2544  0.8370 0.2014
## differential.forest.prediction -0.9320    0.8974 -1.0385 0.8503
```